

PEAS

- Performance Measure
- Environment
- Actuators (output)
- Sensors (input)

Environment

- fully / partially observable
- deterministic / stochastic
 - strategic if det except other agents
- episodic / sequential
- static / dynamic
 - semidynamic if perf score changes but not env
- discrete / continuous
- single / multi-agent

Exploitation vs Exploration Agents

Simple Reflex – just go

Model-based – check pass

Goal-based – check if goal

Utility-based - numbers

Learning agent – any of the above

Problem Formulation

States (state space)

Initial state

Goal states/test

Actions

Transition model

Action cost function

- **Time complexity:** number of nodes expanded
- **Space complexity:** maximum number of nodes in memory
- **Completeness:** return a solution if it exists? (If no solution, does it terminate)
- **Optimality:** always find the least-cost solution?

NB: assume no depth-limiters

Hill Climbing

```
current = initial state
loop:
    neighbor = a highest value successor of current
    if value(neighbor) <= value(current):
        return current
    current = neighbor
```

Simulated Annealing

```
current = initial state
for t = 1 ... ∞:
    T = schedule(t)
    if T = 0:
        return current
    next = a randomly selected successor of current
    if value(next) > value(current) or Prob(next, current, T):
        current = next
```

Local Search

- States (state space)
- Initial state
- Goal test
- Successor function

Minimax

Complete if finite

Time: b^m

Space: bm with dfs

Optimal against optimal

```
def minimax(state):
    v = max_value(state)
    return action in successors(state) with value v

def max_value(state):
    if is_terminal(state): return utility(state)
    v = -∞
    for action, next_state in successors(state):
        v = max(v, min_value(next_state))
    return v

def min_value(state):
    if is_terminal(state): return utility(state)
    v = ∞
    for action, next_state in successors(state):
        v = min(v, max_value(next_state))
    return v
```

```
def minimax_with_cutoff(state):
    v = max_value(state)
    return action in successors(state) with value v

def max_value(state):
    if is_cutoff(state): return eval(state)
    v = -∞
    for action, next_state in successors(state):
        v = max(v, min_value(next_state))
    return v

def min_value(state):
    if is_cutoff(state): return eval(state)
    v = ∞
    for action, next_state in successors(state):
        v = min(v, max_value(next_state))
    return v
```

```
def max_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = -∞
    for action, next_state in successors(state):
        v = max(v, min_value(next_state, α, β))
        if v >= β: return v
        α = max(α, v)
    return v

def min_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = ∞
    for action, next_state in successors(state):
        v = min(v, max_value(next_state, α, β))
        if v <= α: return v
        β = min(β, v)
    return v
```

Search strategy depends on

- Number of goal states
- Distribution of goal states in search tree
- Finite/infinite branching factor/depth
- Repeated states
- Need for optimality?
- Need to know if there is no solution?

Iterative Deepening A* (IDA*)

- Use iterative deepening search
- Cutoff using f-cost [$f(n) = g(n) + h(n)$] instead of depth

Simplified Memory-bounded A* (SMA*)

- Drop the nodes with worst f-cost if memory is full

Heuristic

Admissible: conservative estimate.

Heuristic to relaxed problem is admissible

Theorem: if $h(n)$ is admissible, A* using tree search is optimal

Consistent: every successor n' of n generated by any action a , $h(n) \leq c(n, a, n') + h(n')$

Consistent \rightarrow Admissible

Dominance: if $h_2 n \geq h_1(n)$. Regardless of admissibility

Admissible + A* tree search is optimal

Consistent + A* graph search is optimal

Search	Time	Space	Complete	Optimal	Notes
BFS	b^d	b^d	If finite	If uniform	Store all nodes
DFS	b^m	bm	No if infinite or loops	No	Stack; Space = frontier size
UCS	$b^{C*/\epsilon}$	$b^{C*/\epsilon}$	Yes, if step cost $\geq \epsilon$	Yes	PQ on path cost; Dijkstra; C is cost of optimal solution
DLS	b^l	bl	No	No	
IDS	b^d	bd	If finite	if uniform	0..N
Greedy Best-first	b^m	bm	No esp if bad heuristic	No	$f(n)=h(n)$
A*	b^m	b^m	Yes	Yes*	$f(n)=g(n) + h(n)$

For $b=10, d=5$,

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

Overhead = $(123,456 - 111,111)/111,111 = 11\%$

$\text{Prob}(\text{next}, \text{current}, T) = e^{\frac{\text{value}(\text{next}) - \text{value}(\text{current})}{T}}$

the higher T is, the larger the probability


Theorem: if T decreases slowly enough, simulated annealing will find a global optimum with high probability

Beam Search

Perform k hill-climbing in parallel

Variants:

- Local beam search *pick best nodes*
 - Choose k successors deterministically
- Stochastic beam search *sometimes pick worse*
 - Choose k successors probabilistically



Local vs uninformed search

Search space too large, too long

Solution path not important

Problem with Local optimal

Solubility: the problem configuration is not necessarily solvable, and what we actually want is a configuration that minimizes the number of conflict violations.

$Absolute\ Error = |\hat{y} - y|$
 $Squared\ Error = (\hat{y} - y)^2$

$MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$

$MAE = \frac{1}{N} \sum_{i=1}^N |\hat{y}_i - y_i|$

$Accuracy = \frac{1}{N} \sum_{i=1}^N [\hat{y}_i = y_i]$

Where $\hat{y}_i = h(x_i)$ and $y_i = f(x_i)$.

$Accuracy = \frac{TP+TN}{TP+FN+FP+TN}$

		Actual Label	
Predicted Label	Cancer	2 True Positive	1 False Positive
	Benign	3 False Negative	4 True Negative

$Recall\ R = TP / (TP+FN)$

def DTL(examples, attributes, default):
 if examples is empty: return default
 if examples have the same classification:
 return classification
 if attributes is empty:
 return mode(examples)
 best = choose_attribute(attributes, examples)
 tree = a new decision tree with root best
 for each value v_i of best:
 examples_i = {rows in examples with best = v_i }
 subtree = DTL(examples_i, attributes - best, mode(examples))
 add a branch to tree with label v_i and subtree subtree

Decision tree **consistent** if data can be split on attributes
No. of bool fns
 = no of distinct truth tables with 2^n rows
 = 2^{2^n} distinct decision trees

FP: Type I error
 FN: Type II error

Precision
 $P = TP / (TP+FP)$

F1 Score
 $F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$

$I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$
 $IG(A) = I\left(\frac{p}{p+n}, \frac{n}{p+n}\right) - remainder(A)$

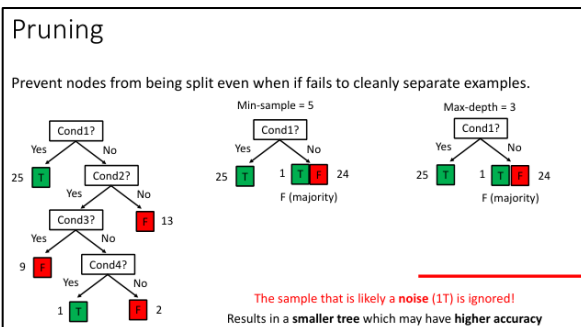
$GainRatio(A) = \frac{IG(A)}{SplitInformation(A)}$
 $SplitInformation(A) = -\sum_{i=1}^d \frac{|E_i|}{|E|} \log_2 \frac{|E_i|}{|E|}$

Make decision trees to use low-cost attributes where possible using Cost-Normalized-Gain:
 $\frac{IG^2(A)}{Cost(A)}$

$\frac{2^{IG(A)} - 1}{(Cost(A) + 1)^w}, w \in [0,1]$

$remainder(A) = \sum_{i=1}^v \frac{|S_i|}{|S|} I(S_i)$

Handwritten notes:
 IDT
 $I() = -\sum p_i \log_2 p_i$
 remainder = $\sum \frac{P_i \cdot M_i}{P \cdot M} I()$
 $IG = I() - R()$
 I(children) weighted to their population
 5 in this side, 12 in this side, 17 in all
 $\frac{5}{17} \left(\frac{1}{5} \log_2 \frac{1}{5} + \frac{4}{5} \log_2 \frac{4}{5} \right) + \frac{12}{17} \left(\frac{6}{12} \log_2 \frac{6}{12} + \frac{6}{12} \log_2 \frac{6}{12} \right)$
 6:7
 CS2109S:



If $h(n)$ is admissible but inconsistent, solution is nonoptimal. Shows graph search is not necessarily better than tree search for A^* , so heuristic must always be chosen carefully.

Adversarial search is only optimal when against optimal opponent. So when dealing with humans, there are much more considerations needed to achieve "optimality".

Making good A^* heuristics is an art in which practice helps.

Multiple ways to handle inconsistent data:

- Pruning: Min-sample and Max-depth.
- Pre-processing of data to remove outliers that create noise.
- Select only relevant features, so less relevant features that could create inconsistencies are not part of the decision tree.
- More data